



CE177


Advanced Operating Systems

Hamid Fadishei, Assistant Professor
Computer Engineering Department,
University of Bojnord
Fall 2018


OS Architectures

Monolithic OS

- All OS services in perivileged mode
- Example: Linux





Linux kernel LOC



Microkernel

Microkernel

- Kernel mode: Minimal services
- The rest runs in user mode
- Philosophy: Not all OS services require privileged access
- Philosophy: Separate policy from mechanism
- Example: CMU Mach, MINIX



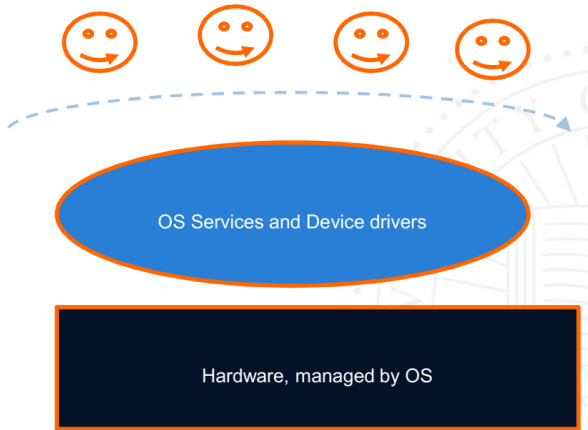


Microkernel benefits

- Microkernels are...
 - Reliable
 - *Because small*
 - *Crash of a component does not necessarily crash the system*
 - Maintainable
 - Extensible
 - *It is easy to add features*

Microkernel drawback

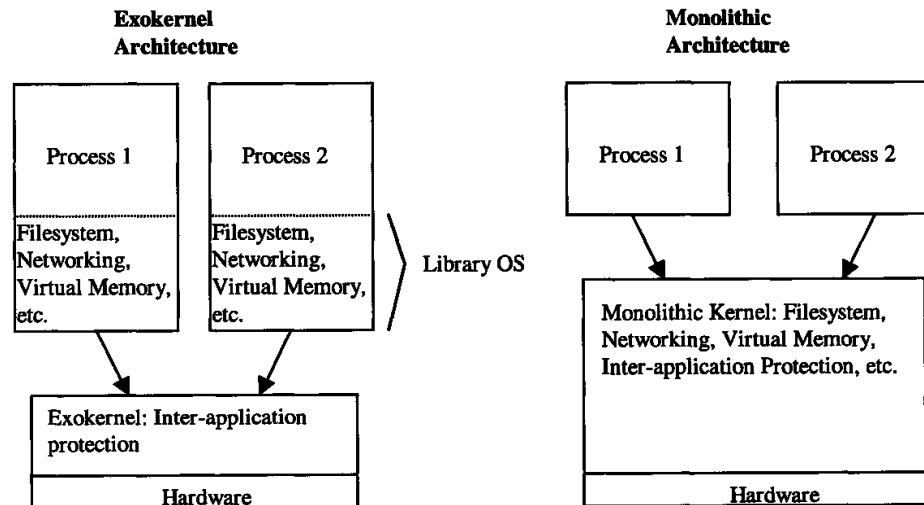
- Performance overhead
 - Cost of crossing the lines
 - *IPC and Syscalls*



Exokernel/Unikernel


Exokernel

- Microkernel and monolithic emphasize on providing an abstraction of hardware to apps
 - Example: read/write services for files (disk)
- Exokernel
 - No emphasis on abstraction at kernel level
 - Developers of specific applications know better than kernel developers how to use the hardware
 - Exokernel takes OS out of kernel and puts it into application (libraries)
 - Separates protection from management




Exokernel

- Exokernel
 - A thin kernel that just perform multiplexing of hardware resources
 - *Memory, TLB, Disk, Network, Frame Buffer, etc*
 - More complex management is inside application lib
 - *FS, TCP, IPC, etc*




Exokernel

- Exokernel systems can be tailored for specific applications
 - High performance
 - Small software footprint
- Example: Cheetah webserver



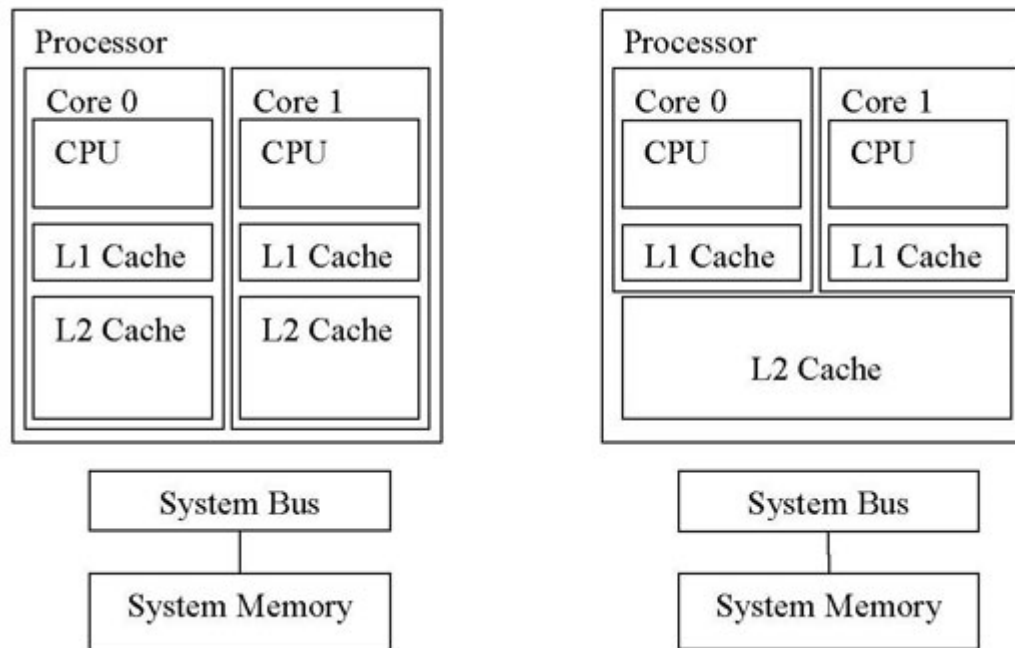
Exokernel Example

- Exokernel example
 - Cheetah web server
- Traditional webserver: 6 buffer operations
 - read() - copy from disk to kernel buffer
 - read() - copy from kernel to user buffer
 - checksum
 - send() - user buffer to kernel buffer
 - send() - kernel buffer to device memory
- Exokernel: 2 operations
 - Copy from disk to memory (packet bodies not files are stored on disk!)
 - checksum
 - Copy from memory to network device

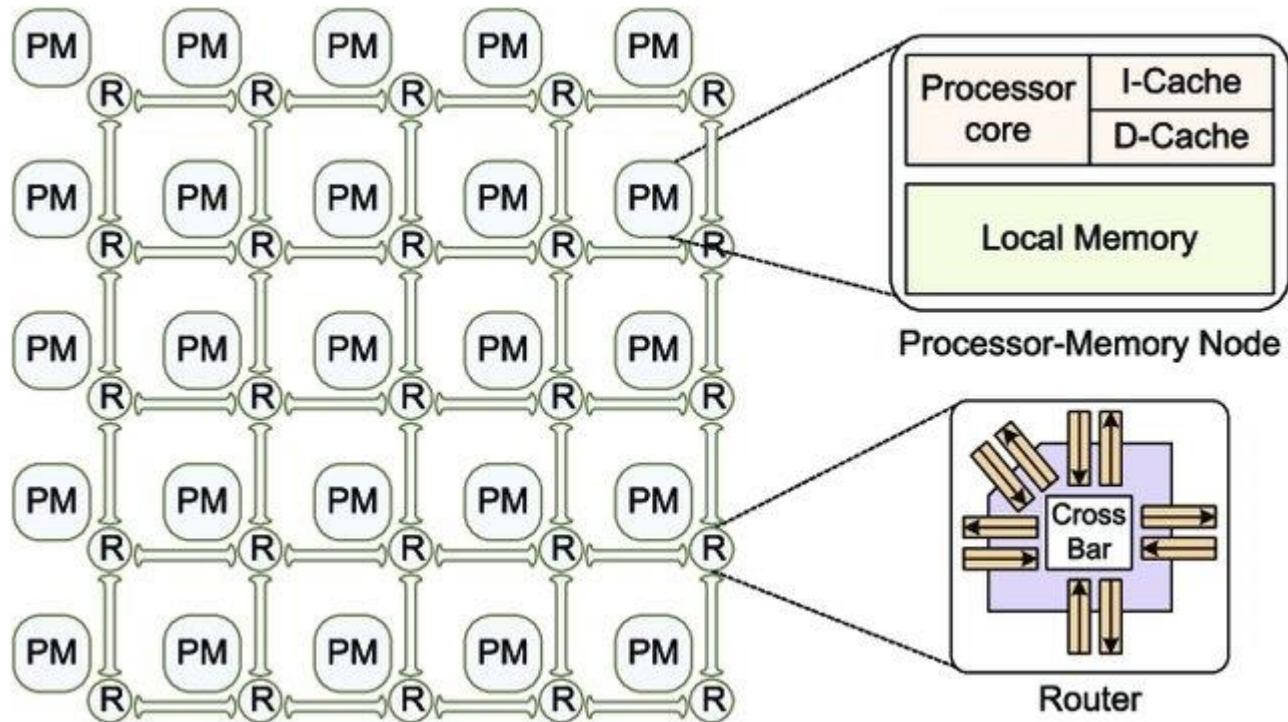


Multikernel

Multicore CPU




Manycore CPU



Multicore/Manycore OS Challenges

- It is difficult to share a data structure among many cores
 - Locks/Atomic instructions => Scalability problem
 - Requires solution at application level
 - *For example, partition data, use lock-free structures, etc*
 - *Requires tremendous engineering*
- Why?
 - Implicit sharing (via shared memory) is troublesome in manycores
 - It is expensive to move cache lines
 - Congestion of interconnects



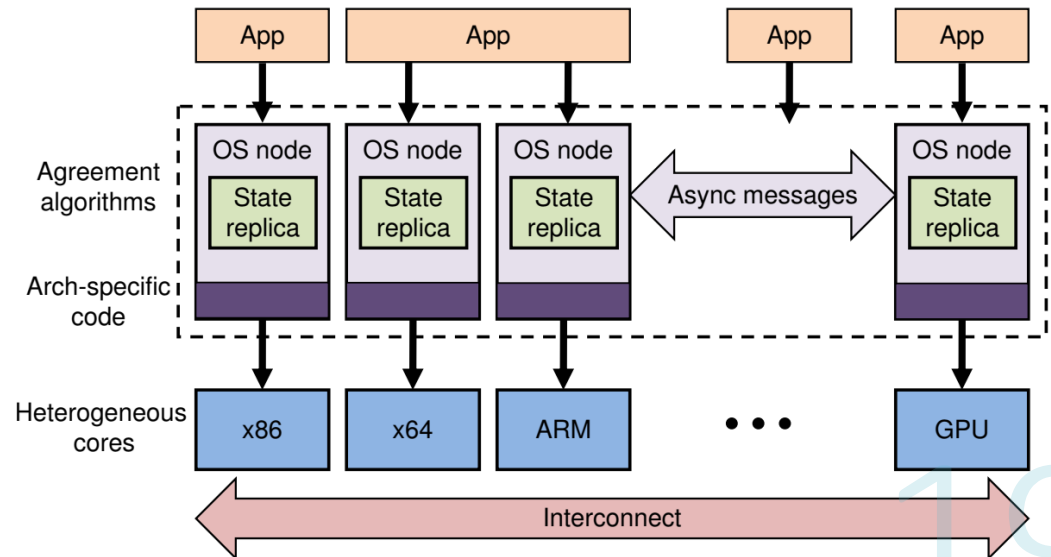
Adopted from OS course by Junfeng Yang <https://www.cs.columbia.edu/~junfeng/12sp-w4118>

Solution: Multikernel

- Multikernel
 - Kernel for multicores
- Multikernel idea
 - Use explicit sharing via messaging (instead of shared memory)
 - No shared data structure
 - Want to see other core's data? Send a message to it!
 - Advantages
 - *Scalable*
 - *Portable (For example for heterogeneous cores)*
 - *Good match for future if new manycores drop support for shared memory coherency*


Challenge: global state

- OS must manage global state among cores
 - Example: Page table of a process
 - If a TLB is modified in a core, all cores should update their TLB
 - AKA *TLB Shootdown*
- Solution (in absence of shared data structure)?
 - Replicate the global state
 - Read => OK, low latency
 - Update =>
 - *Update local copy => OK, low latency*
 - *Distribute the update to other cores => !*



Case study: TLB Shootdown

- Need to keep TLB state of a process consistent among cores
- Conventional method (Linux/Windows)
 - Use shared memory for communications
- Multikernel approach
 - Use message passing for communications



References

- Liedtke, Jochen. *On micro-kernel construction*. Vol. 29. No. 5. ACM, 1995.
- Engler, Dawson R., and M. Frans Kaashoek. *Exokernel: An operating system architecture for application-level resource management*. Vol. 29. No. 5. ACM, 1995.
- Madhavapeddy, Anil, et al. "Unikernels: Library operating systems for the cloud." *Acm Sigplan Notices*. Vol. 48. No. 4. ACM, 2013.
- Baumann, Andrew, et al. "The multikernel: a new OS architecture for scalable multicore systems." *Proceedings of the ACM SIGOPS 22nd symposium on Operating systems principles*. ACM, 2009.